

FA26 — Left vs. Right Fist EEG Classification

What this is

Working notes for the FA26 recruitment challenge (motor imagery/execution, left vs. right fist, PhysioNet EEGMMIDB). This file is the running log + open-questions list the challenge asks for ("log your attempts").

Data, as found

- `data/` = EEGMMIDB, already downloaded: 109 subjects (`S001-S109`), each with 14 runs (`S###R01.edf...S###R14.edf`) + matching `.edf.event` files (ground-truth annotations, redundant with the annotations embedded in the EDF itself — used both as a cross-check).
- Confirmed via `mne.io.read_raw_edf`: 64 EEG channels, nominally 160 Hz, annotations `T0/T1/T2` per run.
- Run mapping (confirmed against the PDF instructions + by checking which runs actually contain T1/T2 vs. only T0):
 - Runs 3, 7, 11: **executed** left/right fist. T1=left, T2=right.
 - Runs 4, 8, 12: **imagined** left/right fist. T1=left, T2=right.
 - Runs 5,9,13 / 6,10,14: both-fists-vs-both-feet (not used for this task).
 - Runs 1, 2: eyes-open/closed baseline (no T1/T2, not used).
- **This task only uses T1/T2 trials, never T0 (rest)** — the assignment is left-fist-vs-right-fist, not fist-vs-rest.

Known data-quality issues

- S088: recorded at 128 Hz instead of 160 Hz, with 19 trials instead of 15, and irregular trial durations (1.38-5.12s vs. a consistent ~4.1s)
- S092: recorded at 128 Hz instead of 160 Hz, with 19 trials instead of 15, and irregular trial durations (1.38-5.12s vs. a consistent ~4.1s)
- S100: recorded at 128 Hz instead of 160 Hz, with 12 trials instead of 15, each abnormally long (~5.1s vs. ~4.1s)
- S089: run duration is 181s instead of the standard 125s, with 22 trials instead of 15
- S104: run duration is 106s instead of the standard 125s, with 13 trials instead of 15
- Trial counts per run are not always exactly 15 (e.g., S034/S037 executed runs have 14 trials instead of 15) — handled naturally since the loader counts actual annotations rather than assuming a fixed count.
- Did **not** exhaustively QC all 109 subjects' impedance/noise levels — see ToDo below. There are documented issues with a handful of EEGMMIDB subjects (channel corruption, mislabeled events) reported by other groups using this dataset; I have not cross-checked against those reports.

Folder Structure

```
data/
|-- ANNOTATORS
|-- RECORDS
```

```
|-- 64_channel_sharbrough.pdf
|-- S001/
|-- |-- S001R01.edf
|-- |-- S001R01.edf.event
|-- S002/ ... (same pattern)
|-- ...
|-- S109/ ... (same pattern)
```

Files

For this project, only Runs # 3, 4, 7, 8, 11, and 12 are used.

Run	Task	Condition	Subjects (.edf files)	.edf.event files
03	Left vs. right fist	Executed	109	109
04	Left vs. right fist	Imagined	109	109
07	Left vs. right fist	Executed	109	109
08	Left vs. right fist	Imagined	109	109
11	Left vs. right fist	Executed	109	109
12	Left vs. right fist	Imagined	109	109

.edf (European Data Format): Actual Signal


- Standard binary file format for storing multi-channel biosignal recordings (EEG, EMG, ECG, etc.).
- Contains: a header (subject metadata, recording start time, number of channels, channel labels, sample rate, physical/digital signal ranges) followed by the raw digitized signal samples (16bit signed integer) for every channel.
- In this dataset: 64 EEG channels sampled at 160 Hz, continuous recording for the full run duration (~2 minutes for runs 3-14).
- Annotations (event markers T1/T2) are also embedded inside the EDF itself as an "EDF+ Annotations" channel

.edf.event: Separate Annotations (e.g., T1: Left Fist, T2: Right Fist)

Data Values

- Digital Value:
 - Each raw digitized signal sample (16bit signed integer) stored in edf file is a measurement value from the EEG electrode
 - An electrode picks up electrical potential from neural activity
 - Amplifier boosts the signal
 - Anti-aliasing Filter is applied to prepare for conversion to digital values
 - ADC converts analog values to digital values sampled at 160Hz
 - Data has been collected from 64 electrodes (channels) as shown in **64_channel_sharbrough.pdf**
 - Each .edf file contains 64 parallel samples captured at 160Hz

- Physical Value:
 - Converted from Digital Value to microvolts (μV) of electrical potential at the scalp
 - Conversion Equation:
 - Visualization of a single channel/electrode (C3) from one subject (S001) during Run #4
 - x axis: time in second
 - y axis: physical value in microvolts μV
 - colored band: Annotation value (T0, T1, T2)



Environment

- Set up a project-local venv at `~/FA26/.venv`
- `requirements.txt` has pinned versions
 - This file lists all the packages that need to be installed for the project to run
- To reproduce, run on Terminal:
 - `cd ~/FA26`
 - `python3.11 -m venv .venv`
 - `.venv/bin/pip install -r requirements.txt`
 - `source .venv/bin/activate`

Pipeline (`src/`)

`data.py`

Function	What it does
<code>_runs_for(run_type)</code>	Maps " <code>executed</code> "/" <code>imagined</code> "/" <code>both</code> " to the corresponding run numbers (3,7,11 / 4,8,12 / all six)
<code>subject_run_path(subject, run, data_dir)</code>	Builds the file path for a given subject+run, e.g. <code>data/S001/S001R04.edf</code>
<code>load_raw(subject, run, data_dir)</code>	Reads one EDF file, standardizes channel names/montage, resamples to 160 Hz if needed
<code>raw_to_epochs(raw, tmin, tmax, l_freq, h_freq)</code>	Bandpass-filters (7-30 Hz) and epochs a loaded <code>raw</code> object into T1/T2 trials
<code>build_dataset(subjects, run_type, ...)</code>	Loops over many subjects/runs, calls <code>load_raw</code> + <code>raw_to_epochs</code> on each, crops to 1-2s window, stacks everything into one <code>Dataset</code>
<code>Dataset</code> (dataclass)	Container for the final <code>X</code> , <code>y</code> , <code>groups</code> , and metadata returned by <code>build_dataset</code>

model.py

Function	What it does
<code>build_pipeline(n_components, reg)</code>	Constructs a scikit-learn Pipeline : CSP (spatial filtering) → LinearDiscriminantAnalysis (classifier)

evaluate.py

Function	What it does
<code>within_subject_cv(build_pipeline_fn, X, y, groups, ...)</code>	Per-subject 5-fold shuffled CV (train/test trials from the same subject) — optimistic upper bound
<code>leave_one_subject_out(build_pipeline_fn, X, y, groups, shuffle_labels, ...)</code>	Trains on N-1 subjects, tests on the held-out one, repeats for all subjects; shuffle_labels=True runs the permutation-baseline control
<code>group_kfold_cv(build_pipeline_fn, X, y, groups, n_splits)</code>	Faster stand-in for LOSO — folds respect subject grouping but hold out several subjects per fold at once
<code>summarize(results, title)</code>	Formats a results dict (mean/std/min/max, confusion matrix) into a printable string

train.py

Function	What it does
<code>parse_subjects(spec)</code>	Parses a CLI string like "1-40" or "1,2,5-10" into a list of subject IDs
<code>main()</code>	Orchestrates the full experiment: parses args → build_dataset() → runs within-subject CV, LOSO, and permutation baseline → writes JSON results → fits and saves a final model (.joblib) on all loaded data

predict.py

Function	What it does
<code>predict_edf(model_path, edf_path)</code>	Loads a saved model + one raw EDF, preprocesses it the same way as training (raw_to_epochs), aligns channels, and returns per-trial predictions (plus true labels if available)
<code>main()</code>	CLI entry point: parses --model/--edf args, calls predict_edf , prints per-trial results and overall accuracy if ground truth is present

plot_results.py

Function	What it does
No functions	Loads <code>results_executed_40.json/results_imagined_40.json</code>
	Builds a two-panel scatter plot of per-subject LOSO accuracy (real vs. permuted labels)
	Saves <code>results/loso_accuracy.png</code>

Roles

- `data.py`: shared toolkit; how to turn a raw EDF into model-ready trials.
- `model.py`: definition of what the model architecture is.
- `evaluate.py`: how to honestly measure accuracy given a dataset and a trained model.
- `train.py`: run the whole training + evaluation experiment end-to-end.
- `predict.py`: test the trained model. it classifies one new file with an already-trained model.
- `plot_results.py`: turn saved results from training + evaluation into a figure (not from predict.py)

How to Run

1. Run Training

- `.venv/bin/python src/train.py --subjects 1-40 --run_type executed --tag executed_40`
- `.venv/bin/python src/train.py --subjects 1-40 --run_type imagined --tag imagined_40`
- `.venv/bin/python src/train.py --subjects 1-109 --run_type executed --loso_only --tag executed_full`
- The training dataset can be selected using the `--subjects` flag
 - To select non-consecutive subjects for training, use **`--subjects 1-10, 33, 50-55`**
- To remove the 5 bad datasets from training, set the `--exclude_bad` flag, for example,
 - `.venv/bin/python src/train.py --subjects 1-40 --run_type imagined --exclude_bad --tag imagined_40`
- During training, the data (selected by the `--subjects` flag) is divided into 2 sets, training set and evaluation set.
 - The training set is used to train the model.
 - The evaluation set is used to evaluate the accuracy of the trained model.
 - We want the evaluation set to be different from the training set so that we don't cheat (test on the training set)
 - We can dictate how much data is used for training and how much for evaluation using the `--train_ratios` flag.
 - for example, if you set `--train_ratios 0.8`, then 80% of the data will be used for training and 20% for evaluation
 - Run **`.venv/bin/python src/train.py --subjects 1-40 --run_type imagined --exclude_bad --train_ratios 0.8 --tag imagined_40`**

- You can train the model with both "executed" and "imagined" signals **--run_type both**
 - This will include all of 6 runs (3 executed and 3 imagined) from each subject
 - This design does not allow you to choose from the 6 runs in each subject
 - To test different combinations of training sets between executed and imagined, you set --run_type to both and then select list of subjects to train with
- Training produces two files
 - A JSON File: /results/results_.json
 - Trained Model: /results/model_.joblib

2. Run Inference (Prediction)

- After your model is trained, you can test it by running it with a new input data
- Your trained model from Step #1 above is a joblib file (e.g., model_executed_40.joblib)
- Your test input data is an edf file (e.g., S050R03.edf)
- Run: `.venv/bin/python src/predict.py --model results/model_executed_40.joblib --edf data/S050/S050R03.edf`
- Prediction prints out the classification result: left_fist or right_fist
- Subjects that were not included in the training session via --subjects can be used for inference tests
 - These are not used for model training, so there is no over-fitting (it's cheating to train and test on the same data)
 - For example, if you ran **train.py --subjects 1-89**, then all subjects 90-109 are free to use for inference

3. Visualize Trained Model

- Visualize accuracy evaluation result of the trained model from
 - results/results_executed_40.json
 - results/results_imagined_40.json
- Run **.venv/bin/python src/plot_results.py**

Evaluation Design (why 3 regimes, not one number)

1. **Within-subject CV** (**ShuffleSplit**, 80/20, 5 repeats, per subject): most optimistic number. CSP/LDA get to fit that person's exact head geometry/noise. Useful as an upper bound, not as evidence of generalization.
2. **Leave-one-subject-out (LOSO)**: fit on N-1 subjects, test on the held out one, repeated for all subjects. This is the number that matters for "does this work on a new person" - the framing the prompt cares about.
3. **Permutation baseline**: same LOSO procedure, but labels shuffled *within each subject* before fitting (preserves per-subject class balance, breaks the left/right correspondence). If real is close to permuted, the real number isn't measuring left-vs-right content.

All subject-grouped splits use subject IDs as groups so no subject's trials appear in both train and test in the same fold - the single most important control here, since within-subject leakage (same head, same session noise) is confounded with the fist label unless explicitly split out.

Results so far (n=40 subjects, S001-S040)

Regime	Executed (3,7,11)	Imagined (4,8,12)
Within-subject CV	0.579 ($\pm$ 0.122)	0.536 ($\pm$ 0.164)
LOSO, real labels	0.543 ($\pm$ 0.070)	0.546 ($\pm$ 0.092)
LOSO, permuted labels	0.478 ($\pm$ 0.048)	0.503 ($\pm$ 0.049)
Paired t-test (real vs. permuted, per-subject)	t=4.26, p=0.0001	t=2.85, p=0.007

See [results/loso_accuracy.png](#) for the full per-subject spread (sorted),
[results/results_executed_40.json](#) / [results/results_imagined_40.json](#) for raw numbers,
[results/model_executed_40.joblib](#) / [results/model_imagined_40.joblib](#) for the saved pipelines
predict.py loads.

Reading of these numbers:

- LOSO real accuracy (~0.54-0.55) is statistically distinguishable from the permutation baseline (~0.48-0.50) — there is a real, if weak, cross-subject left/right signal being picked up, consistent with the literature's characterization of this task as "weak, noisy, unreliable."
- Executed and imagined are statistically similar at this scale (both ~0.54), which was mildly surprising — I expected executed to show a clearer LOSO advantage given its stronger single-subject signal (within-subject CV is a bit higher for executed: 0.579 vs 0.536). More subjects and/or a paired same-subject executed-vs-imagined comparison would sharpen this.
- Per-subject spread is wide (min 0.467/0.378, max 0.778/0.822) — a single mean accuracy badly understates how variable this is. Some subjects are near chance under real labels; a few are much better than the mean. This is exactly the "range, not best case" the prompt asks for.
- Within-subject CV overstates generalizable performance relative to LOSO, as expected, but the gap is smaller than I initially guessed — suggesting the CSP+LDA model isn't wildly overfitting to individual subject idiosyncrasies once shrinkage-regularized; it's just that the cross-subject signal itself is weak.

What I have NOT yet done (ToDo)

- ☐ Experiment with different number of subjects to train, different values of run_type
- ☐ Scale LOSO to the full 109 subjects (currently n=40, chosen for turnaround time in this pass — training + evaluating all 109 with full LOSO is ~3x this runtime; queued as a follow-up run).
- ☐ Quantify "how much of this is about the person, not the task" more directly: e.g., train a subject-identity classifier on the same features and see how well subject ID alone is decodable, to bound how much of any left/right signal could ride on subject-correlated artifacts. Currently I only have the permutation-baseline control, which addresses label-shuffling but not a full subject-confound audit (e.g., electrode impedance differences, session-to-session amplifier drift).
- ☐ Test whether a model trained on **executed** runs transfers to **imagined** runs (and vice versa) — explicitly suggested in the prompt as "fair game and possibly interesting." Not yet run.

- ☐ Try a channel subset restricted to sensorimotor electrodes (C3, C4, Cz and neighbors) instead of all 64 — both as a sanity check (does CSP concentrate its spatial patterns there, as physiology predicts?) and as a way to test robustness to electrode-count mismatches on other EDF files.
- ☐ Look at what CSP is actually learning: plot the spatial patterns (`CSP.plot_patterns`) for a few subjects to check whether the topography is plausible (contralateral sensorimotor, roughly symmetric between left/right) rather than picking up an artifact (e.g., a bad channel, eye-movement contamination near frontal electrodes).
- ☐ Sweep the crop window (currently fixed 1-2s post-cue) and the band-pass range (currently fixed 7-30 Hz) — both were taken from the standard MNE tutorial defaults referenced in the challenge PDF, not independently tuned or justified for this dataset.
- ☐ Cross-check accuracy against a hand-crafted spectral-power baseline (log-bandpower in mu/beta per channel -> LDA, no CSP) to isolate how much CSP's spatial filtering specifically contributes versus a simpler feature.
- ☐ Full data QC pass: I have not screened all 109 subjects for bad channels / flat channels / motion artifact beyond the coarse sfreq/trial-count check already done for S088/S092/S100. Some literature on EEGMMIDB reports specific subjects with known recording issues — I have not cross-referenced that list.
- ☐ `predict.py` currently assumes the EDF has all 64 standard channels in a nameable/standardizable form (via `mne.datasets.eegbci.standardize`). Have not tested it against a file with a different channel montage or a partial channel set — worth hardening before calling it a real deliverable-grade CLI.
- ☐ Demo video, README with pinned one-line usage example, and GitHub repo publication are not done yet — this session focused on the technical pipeline and its evaluation.

Open questions / things I'd want to defend or am least sure of

- Is 1-2s post-cue the right window for *executed* trials specifically? Movement onset latency for executed fist clenches may differ from imagined ones; using the same fixed crop for both may quietly disadvantage one condition. Not yet checked.
- The permutation baseline shuffles within-subject, which controls for subject-level confounds but not for run-order or fatigue/drift-within-run confounds (e.g., if T1 trials cluster earlier in a run than T2 for incidental protocol reasons, permuting within-subject-across-all-trials won't catch a within-run temporal bias). Would want to also try shuffling within-run rather than within-subject to rule this out.
- LOPO with only 40 subjects gives per-subject accuracy estimates with meaningful sampling noise (as few as ~42-45 trials per held-out subject for executed runs). The between-subject std (~0.07-0.09) partly reflects subject differences and partly reflects small-n measurement noise per subject — haven't decomposed those two sources.
- Haven't yet checked whether accuracy correlates with anything measurable about a subject (session length, recording date, channel noise level) that would indicate a data-quality confound rather than a genuine physiological difference in imagery/execution ability.

AI tool use

Used this AI assistant (Enchanté, running Claude) to read the challenge PDFs, inspect the dataset structure interactively (EDF header content, run annotation labels, sfreq irregularities), scaffold the pipeline code above, and run/summarize experiments. One specific choice made and reasoned about

directly (not just accepted from the tutorial default): using Ledoit-Wolf-shrinkage-regularized CSP covariance (`reg='ledoit_wolf'`) rather than the tutorial's unregularized default (`reg=None`) — the 64-channel spatial covariance is estimated from only ~160 time samples per trial after cropping to the 1s analysis window, which is close to rank-deficient (fewer samples than channels means the raw sample covariance is not even guaranteed full rank), so shrinkage was chosen deliberately rather than left at the tutorial's default.